

# COMP417 - Intro to Robotics

Optimal Control

Fall 2019

David Meger



McGill

**MRL** Mobile Robotics Lab  
at **McGill University**


# Intro to Optimal Control

- Rather than having a reference trajectory, suppose we have a cost function that tells us just how bad each “wrong” action is
  - Achieving goal can still be the best (least cost)
  - Now we can prioritize different mistakes
- This allows us to ***think ahead in time*** and solve for trajectories that minimize the sum of costs as we formulate their matching control

# The state of a double integrator (block on ice)

$$\mathbf{x} = \begin{bmatrix} G \\ p_x \end{bmatrix}$$


State = [Position along x-axis]



# Controls of a double integrator

$$\mathbf{u} = \begin{bmatrix} G \\ u_x \end{bmatrix}$$


Controls = [Force along x-axis]



# Dynamics of a double integrator

Passive dynamics correspond to taxi sliding on the roads of Montreal (dirty ice). Where is the car going to end up?  
Similar to curling.

$$F = ma$$
$$-k\dot{x} = m\ddot{x}$$



# Dynamics of a double integrator

Passive dynamics correspond to taxi sliding on the roads of Montreal (dirty ice). Where is the car going to end up?  
Similar to curling.

$$F = ma$$

$$-k\dot{x} = m\ddot{x}$$


Controlled dynamics:

$$m\ddot{x} = u - k\dot{x}$$

What controller achieves  $\ddot{x}_{ref}$ ?

$$u_{opt} = \frac{\ddot{x}_{ref} + k\dot{x}}{m}$$


If we assume a known  
reference acceleration, this is enough.  
Let's add some challenge!



# Dynamics of a double integrator

Controlled dynamics interesting when force is limited.


$$\ddot{x} = F - k\dot{x}$$
$$\text{s.t. } |F| < u_{max}$$



Can we design controllers that achieves  $x_{goal} = 0$ ?

- As quickly as possible (get to your meeting on time)
- Without over-shooting (suppose the taxi is parking at your front door!)
- Using as little energy as possible (gas costs)

# Phase plots: A useful analytical tool



Graph how the system evolves in terms of  $x$  and  $\dot{x}$

- This lets us start to reason about what's possible and hints at the ideas we'll use for important control algorithms

For the idealized block on ice (no friction for a moment):

$$\ddot{x} = u, \text{ let } u = -1$$

$$\dot{x}(t) = \dot{x}(0) - t$$

$$x(t) = x(0) + t\dot{x}(0) - \frac{1}{2}t^2$$


Solve for  $t$  in  $\dot{x}$  equation:  $t = \dot{x}(0) - \dot{x}(t)$

Substitute into  $x$  equation

$$x(t) = x(0) + (\dot{x}(0) - \dot{x}(t))\dot{x}(0) - \frac{1}{2}(\dot{x}(0) - \dot{x}(t))^2$$

Note this is quadratic in  $\dot{x}$ , linear in  $x$ . Form:

$$\dot{x}(t)^2 = a x(t) + b$$





Phase Plot:  $u=0$

Phase Plot:  $u=+1$

Phase Plot:  $u=-1$

# An Intuitive Result

- Claim: Bang-bang controller is optimal for the shortest time problem
  - Why?

# An Intuitive Result

- Claim: Bang-bang controller is optimal for the shortest time problem
  - Why?
  - Consider any controller that “pushes” less hard at first. Think about its phase portrait. This will, on average, be slower than the bang-bang. Not optimal.

# An Intuitive Result

- Claim: Bang-bang controller is optimal for the shortest time problem
  - Why?
  - Consider any controller that “pushes” less hard at first. Think about its phase portrait. This will, on average, be slower than the bang-bang. Not optimal.
  - Consider any controller that “brakes” less hard at the critical point. It will overshoot the goal and have to come back around. This will have a faster velocity sometimes, but take longer overall
    - (harder to prove precisely, but you can do it as an exercise)

# Optimal Control

- Formulate control problem as optimization of a cost function given some form of knowledge about the system

$$\mathbf{x}_{t+1} = \mathbf{f}(\mathbf{x}, \mathbf{u})$$

$$g(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{x}_t^T Q \mathbf{x}_t + \mathbf{u}_t^T R \mathbf{u}_t$$

$$\operatorname{argmin}_{u_0, \dots, u_N} \sum_{t=0}^{t=N} g(\mathbf{x}_t, \mathbf{u}_t)$$

- Ideal solution for a robot is over continuous  $\mathbf{x}$ ,  $\mathbf{u}$ , but we could discretize those to make some progress first off...


# Dynamic Programming for Control

- Dynamic Programming (value iteration)
  - Discretize state-space, time and controls
  - Form discrete transition matrix that approximates dynamics,  $f$
  - Form discrete cost that approximates cost,  $g$
- Initialize and update optimal value function (cost-to-go) over states:
  - $J^*(x) = \min_u [g(x, u) + J^*(f(x, u))]$
- Important claims:
  - Knowing  $J^*$  everywhere gives us the optimal controller
$$u^* = \underset{u}{\operatorname{argmin}} [g(x, u) + J^*(f(x, u))]$$
  - Knowledge of  $J^*$  in next states makes it easy to improve  $J^*$  locally, and a simple algorithm that makes an initial guess and updates everywhere converges! (Value Iteration)




# Double Integrator Example


- Goal: arrive at  $x=0$  as soon as possible
- Control  $a(t)=u$ , limited to  $|u| \leq 1$
- Ideally solve over all  $x(0), v(0)$
- Dynamics:
  - $v(t) = v(0) + ut$
  - $x(t) = x(0) + v(0)t + 0.5t^2$
- Cost (min-time):
  - $g(x,u) = 0$  if goal, else 1
- What is the intuitive solution?
- What does dynamic programming look like?




# Value Iteration for the Double Integrator




# Value Iteration for the Double Integrator




# Value Iteration for the Double Integrator



# Value Iteration for the Double Integrator




# Value Iteration for the Double Integrator



# Issues with dynamic programming

- Errors in formulation are related to "size" of discretization
- Naïve discretization scales exponentially with # of dimensions
- Better ways to discretize:
  - Multi-resolution and adaptive grids
  - Trajectory-based methods
- Improving robustness of solution to local errors:
  - 2nd derivative methods
  - Line search
  - Guiding prior information

# State and control of a cartpole



$$\mathbf{x} = [{}^G p_x, {}^G \dot{p}_x, {}^G \theta, {}^G \dot{\theta}]$$

State = [Position and velocity of cart, orientation  
and angular velocity of pole]

$$\mathbf{u} = [f]$$

Control = [Horizontal force]



# Cartpole properties

- Theta joint lacks a motor making this system underactuated
- We must sometimes sacrifice desirable cart position in order to "catch" the pole and right it
- This coupling comes from the dynamics equations
- Two canonical tasks:
  - Swing-up
  - Balancing

